

HIGH QUALITY TELEPHONY USING A FAIL-SAFE MEDIA RELAY SETUP



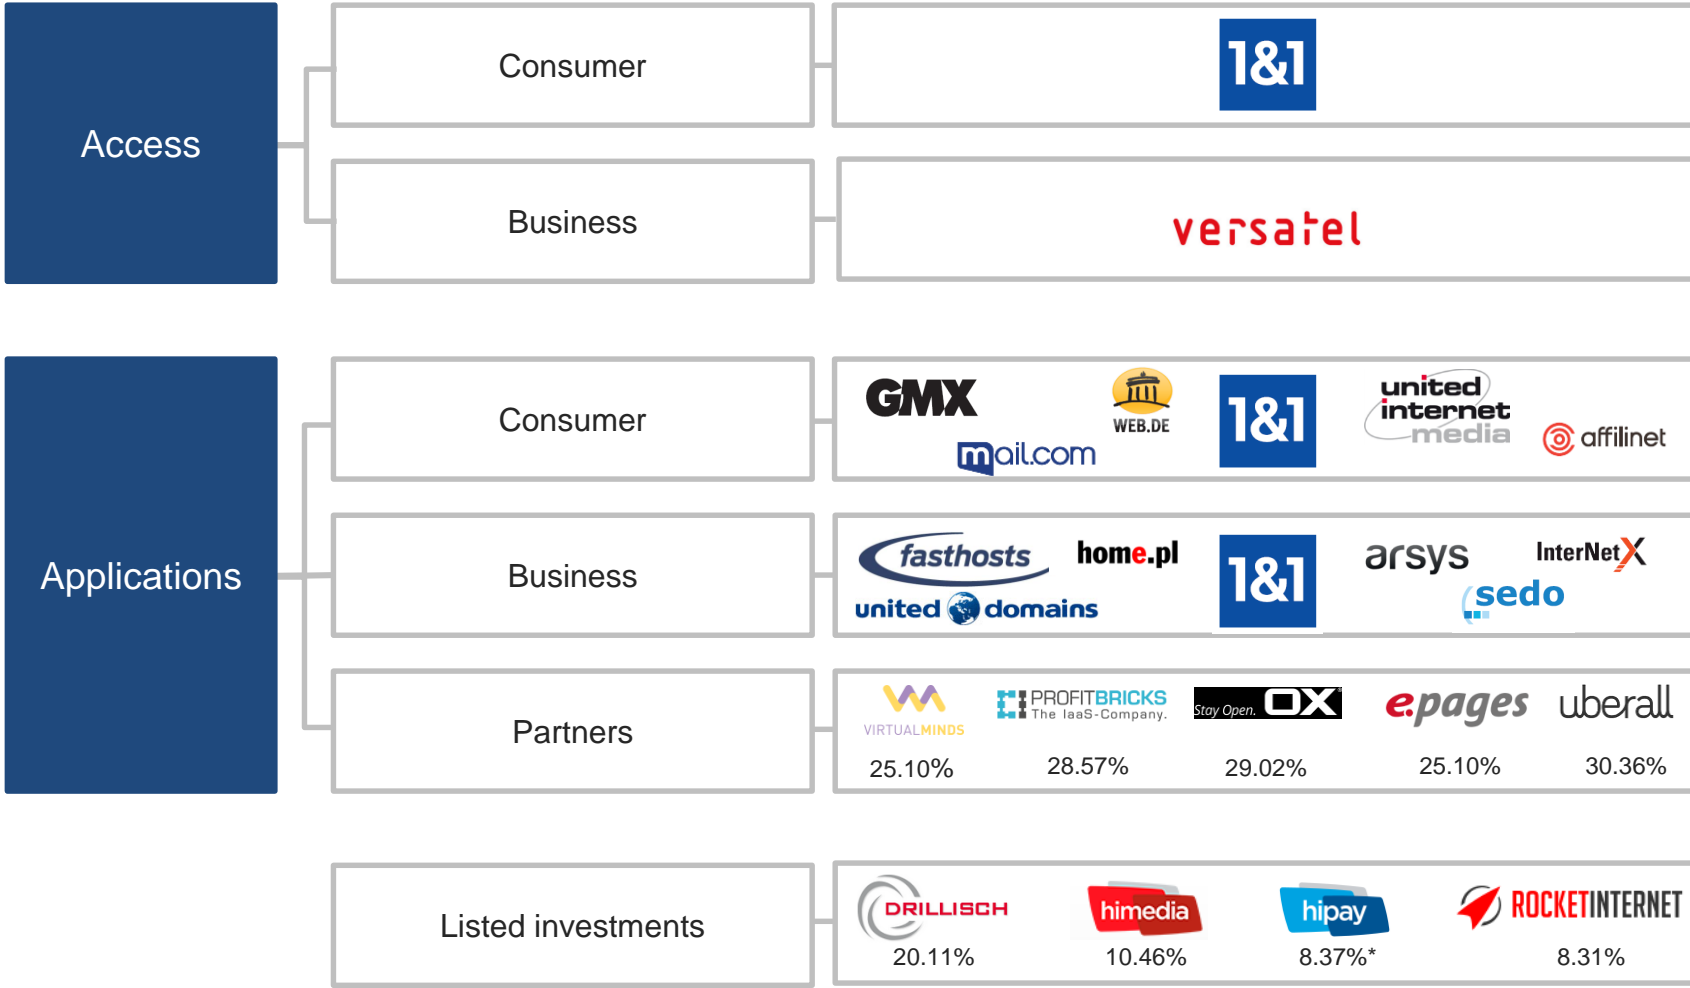
Pawel Kuzak
VoIP Backend Development

Agenda

- 1&1 - Member of United Internet AG
- VoIP Network Overview
 - ☐ Quality of Service
- Media relay - RTPengine
 - ☐ Extended VoIP Architecture
 - ☐ Hardware and Network configuration
 - ☐ Challenges
- Redundancy
 - ☐ RTPengine setup
 - ☐ Redis Keyspace Notifications
- Roundup

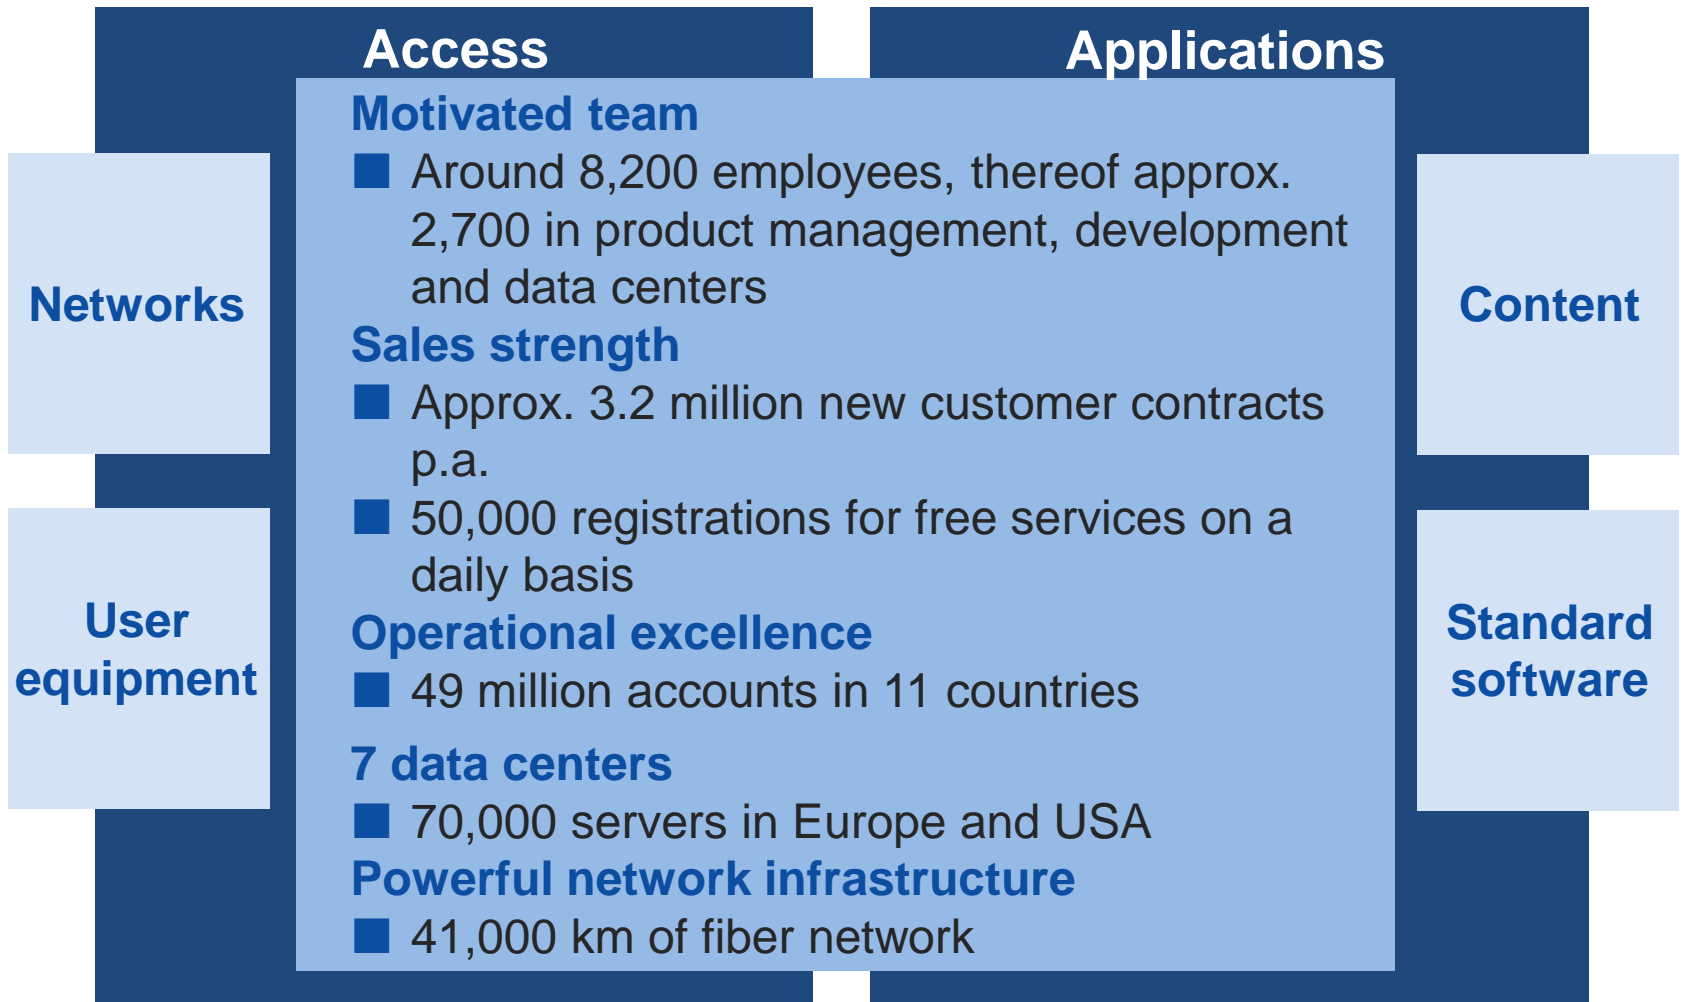


1&1 – Member of United Internet AG



* Spin-off of Hi-Media S.A.

1&1: Internet services of United Internet AG



Locations

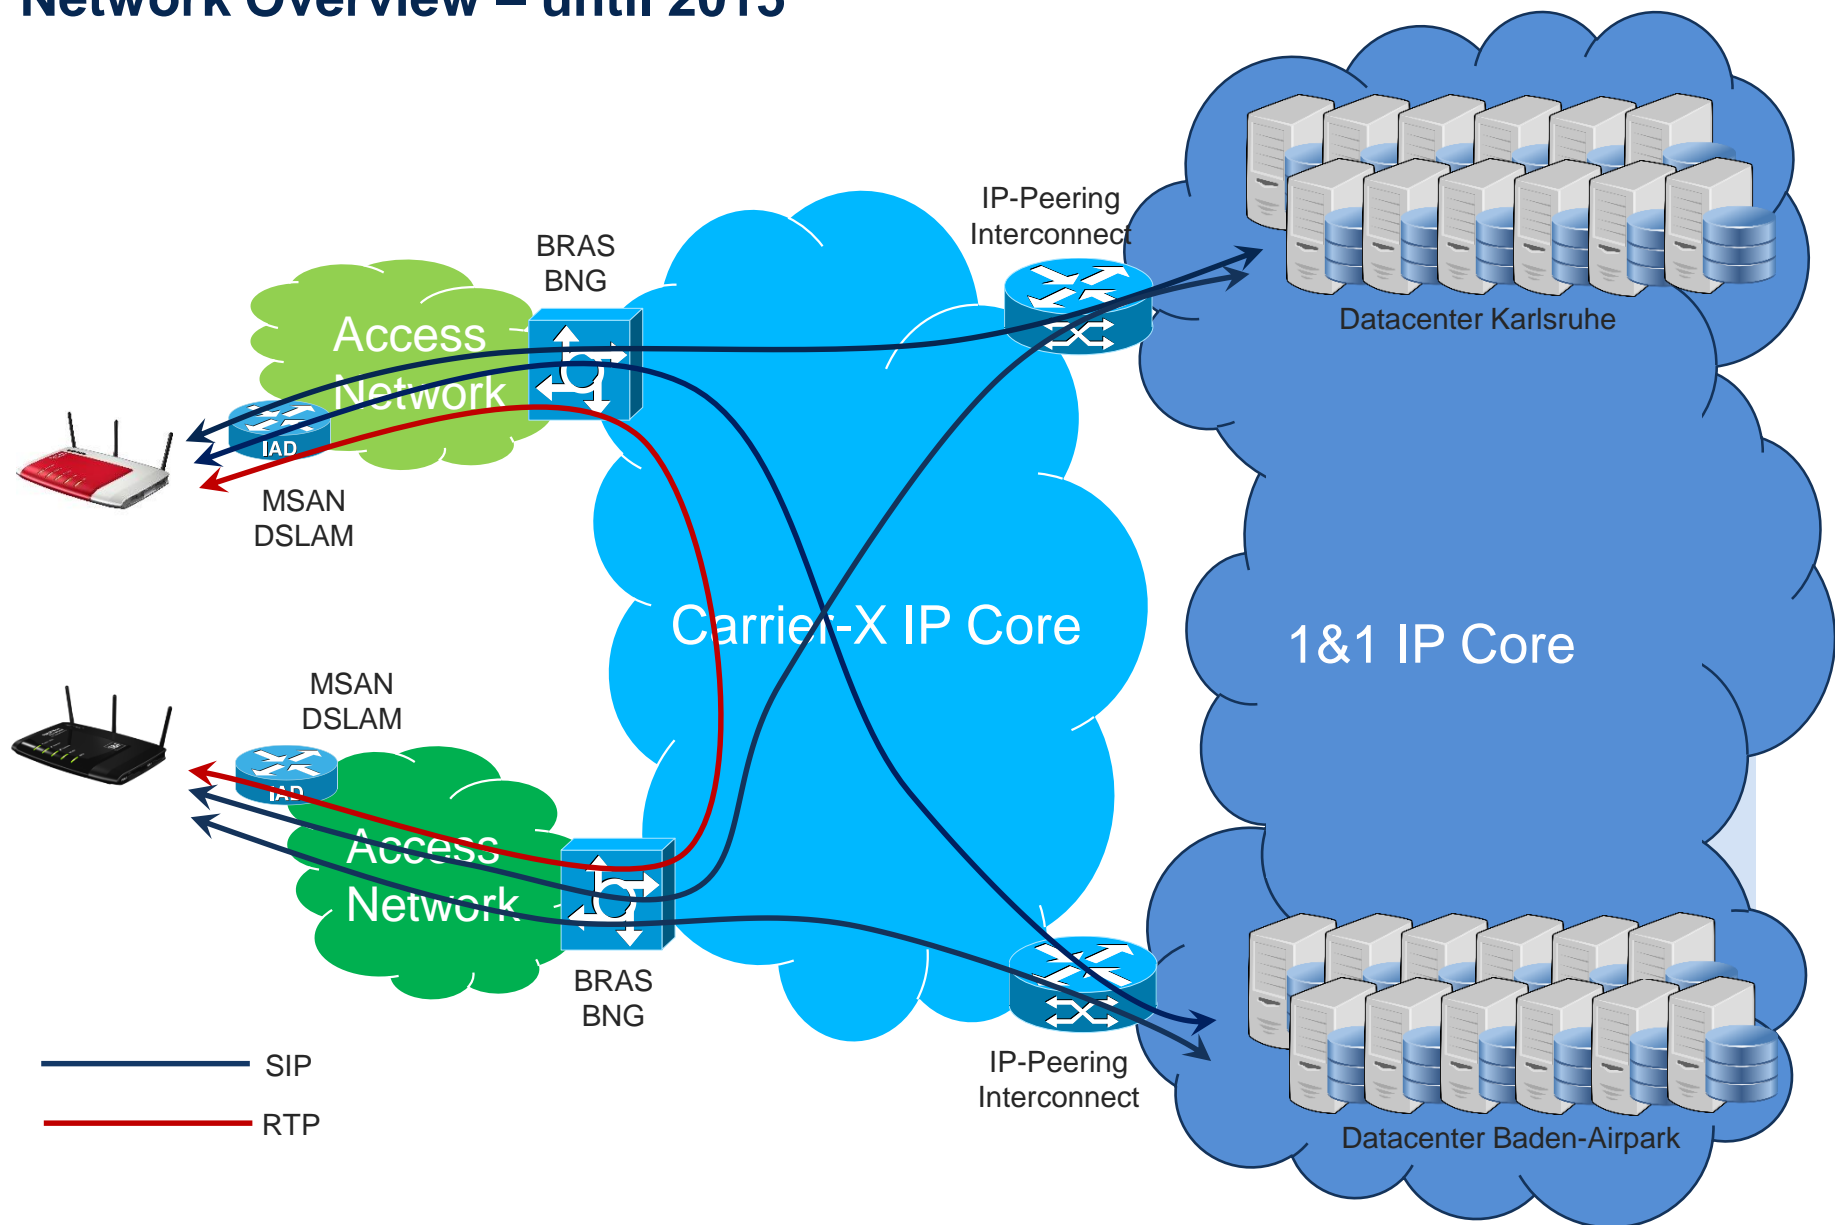


VoIP Backend – Some numbers

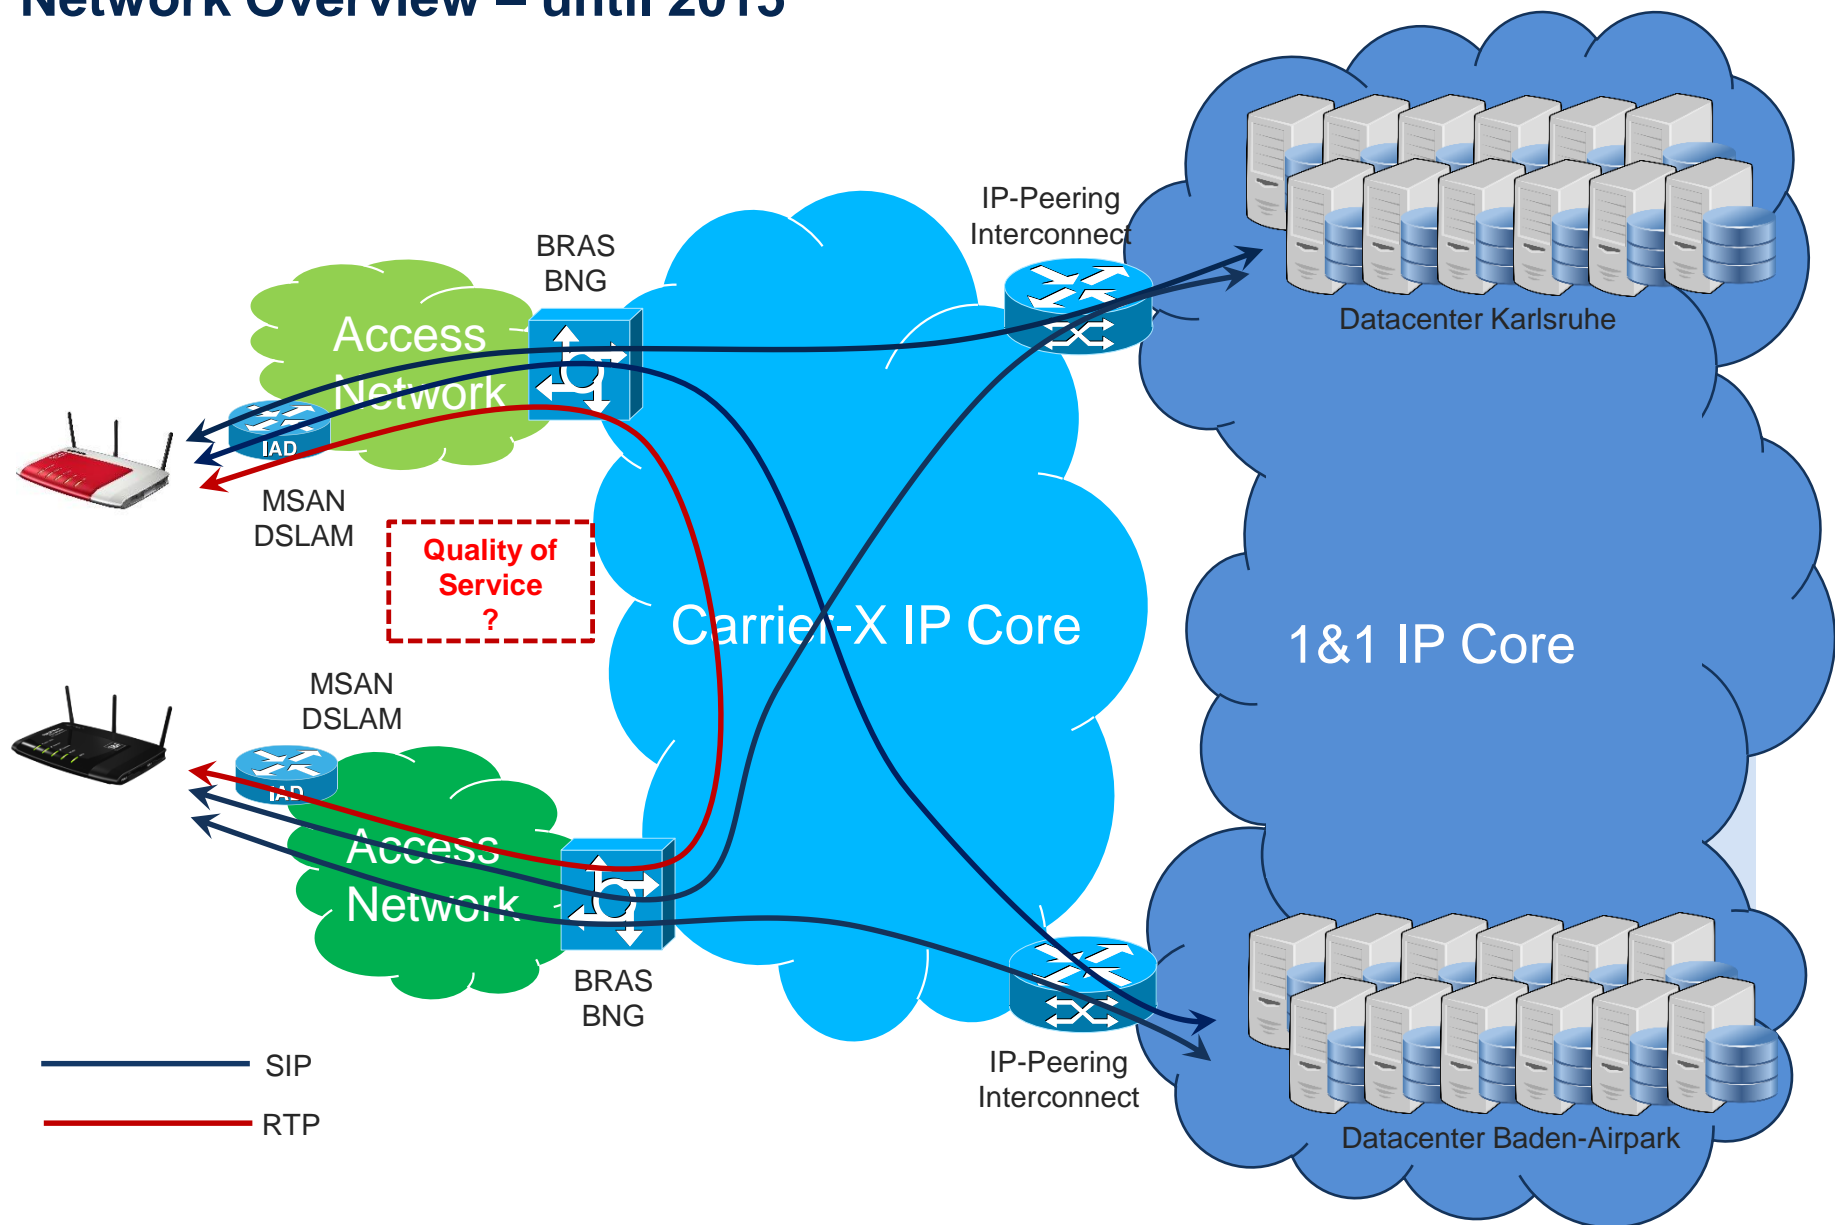
- > 100 servers forming a geo-redundant active-active setup
- ~ 4 Mio. VoIP customers
- ~ 12 Mio. registered numbers
- > 500 Calls/Second (Peak)
- > 120.000 concurrent calls (Peak)
- > 15.000 REGISTER requests per Second (Peak)
- 99,999% availability
- SIP interconnections with
 - Versatel
 - QSC
 - Telefonica
 - Vodafone



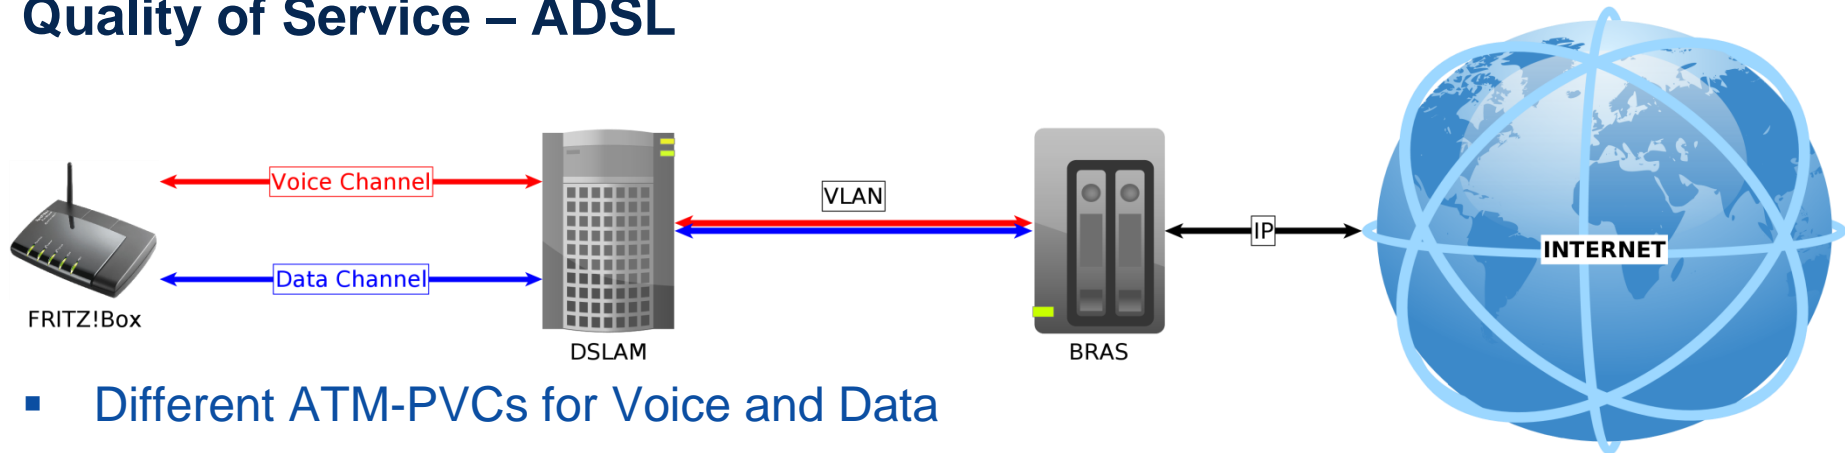
Network Overview – until 2015



Network Overview – until 2015



Quality of Service – ADSL



- Different ATM-PVCs for Voice and Data
- One PPPoE session per ATM-PVC
- Different VLANs for Voice and Data in Backbone


Upstream (F!B → Internet):

- F!B just separates Data traffic from Voice traffic
- Based on the ATM-PVC, the DSLAM decides whether to use the prioritized VLAN to BRAS
- BRAS uses DSCP (sets TOS Bit)

Downstream (Internet → F!B):

- Based on the target IP, the BRAS knows which VLAN to use
- Traffic may be routed best-effort to BRAS

Quality of Service – VDSL



- No distinction between Voice and Data
- One VLAN for Voice and Data in Backbone

Quality of Service – VDSL


Carrier won't prioritize IP traffic unless it doesn't come from a pre-defined IP range

- ➔ Prioritization of end-to-end RTP traffic not possible
 - ➔ IPs spread
- ➔ Need to bundle RTP traffic
- ➔ Need media relay

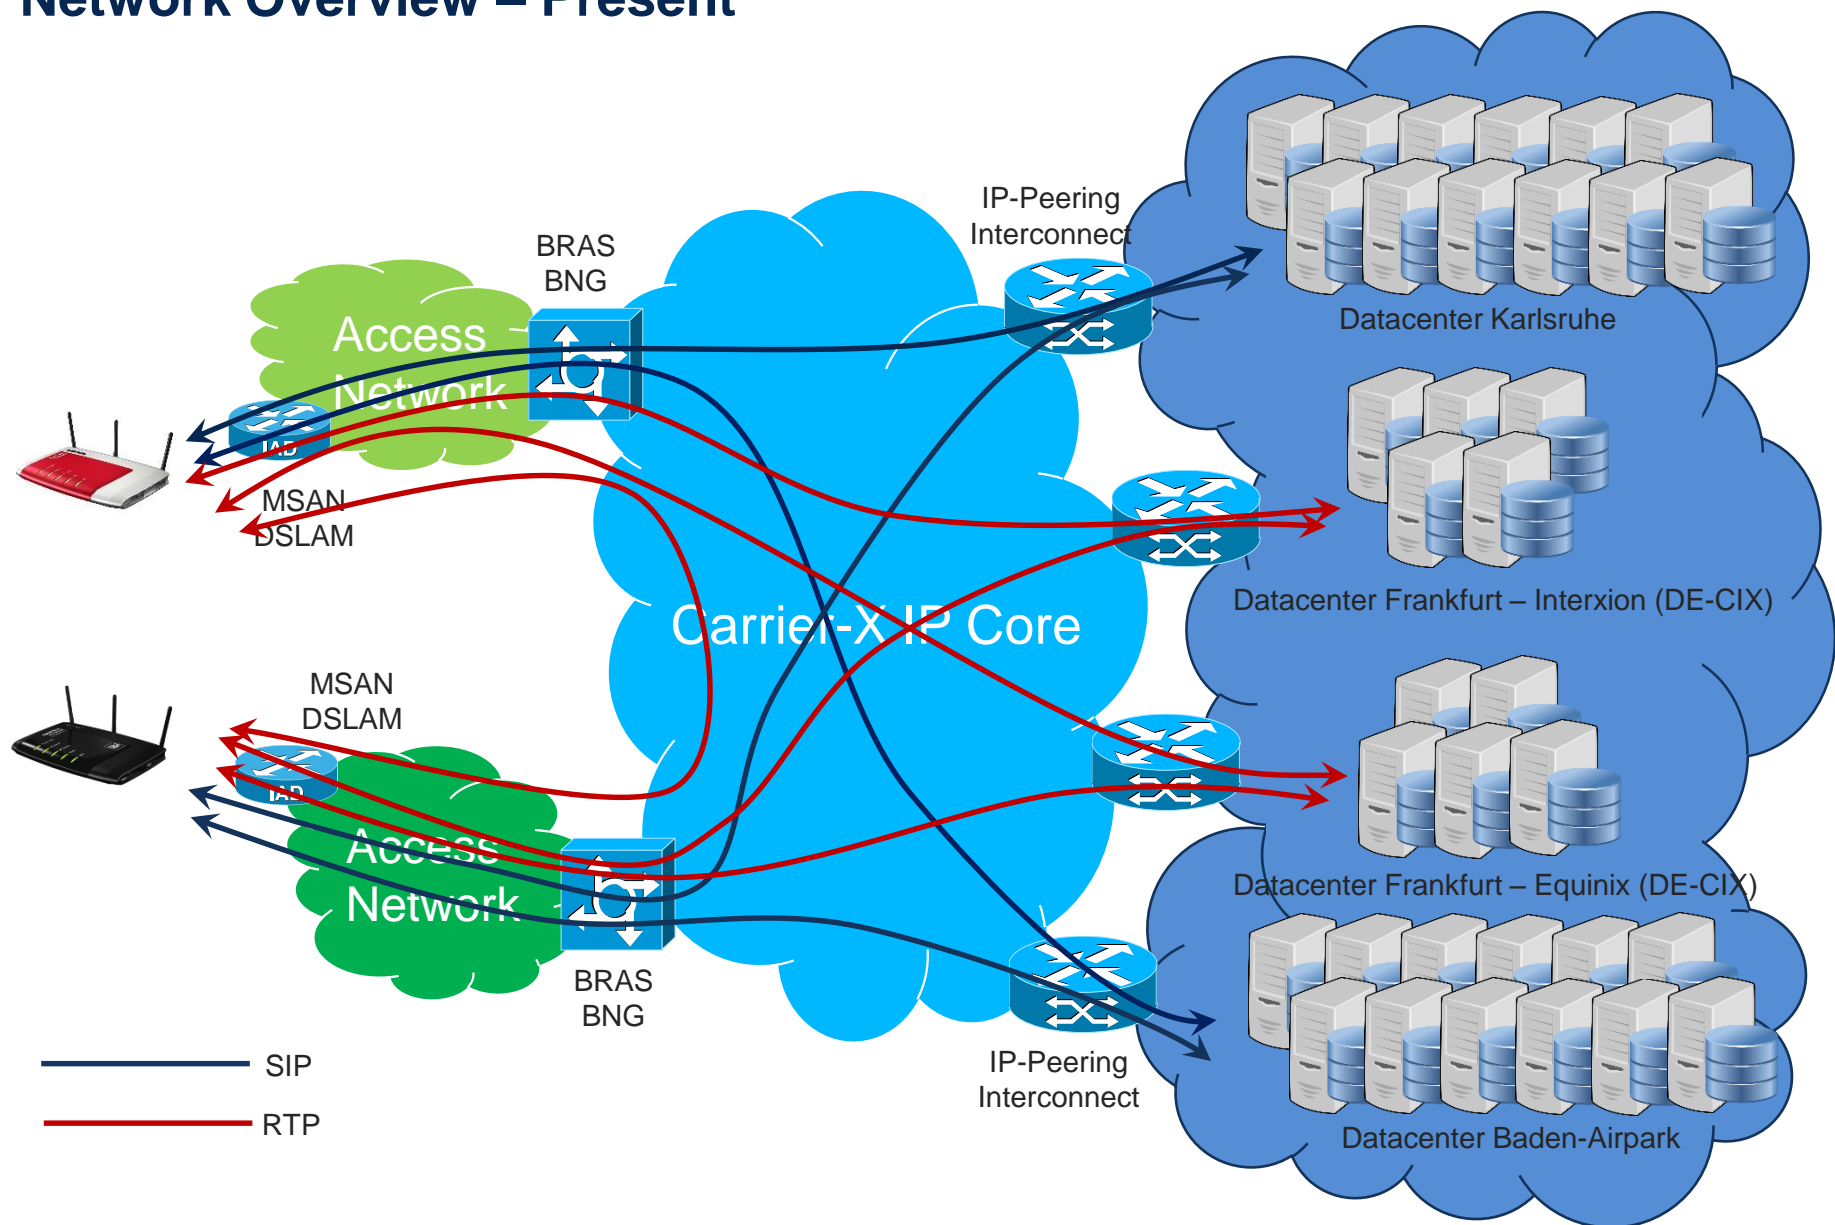
Requirements for media relay:

- Open source
- Kamailio integration
- High performance
- High availability
- Scalability
- TOS – Bit setting


➔ RTPengine (<https://github.com/sipwise/rtpengine>)



Network Overview – Present



QoS impact – VDSL: Speech delay VDSL-ISDN



Media relay – Some numbers

- Currently 6 server in 2 data centers
 - Dell R430, Intel X520 Network 10 Gbit/s, 2x 10 Core Xeon E5 w/ HT
 - Running RTPengine controlled via Kamailio's rtpengine module
- Full redundant active-active setup
- Each server handling 3.000 concurrent sessions (Peak)
 - ~ 300.000 packets/s incoming and outgoing
- Each server able to handle 18.000 concurrent sessions (Peak)
 - Worst case backup scenario
 - ~ 18.000.000 packets/s incoming and outgoing
- Successfully tested 25.000 concurrent sessions

Limiting factors:


- Bandwidth → Extend hardware
- Available sockets → Use more IPs
- Interrupts → Tune configuration



Media relay – Interrupt handling configuration

Intel® Flow director / N-tuple filter / Hash filter:
→ Map RTP/RTCP to same queue

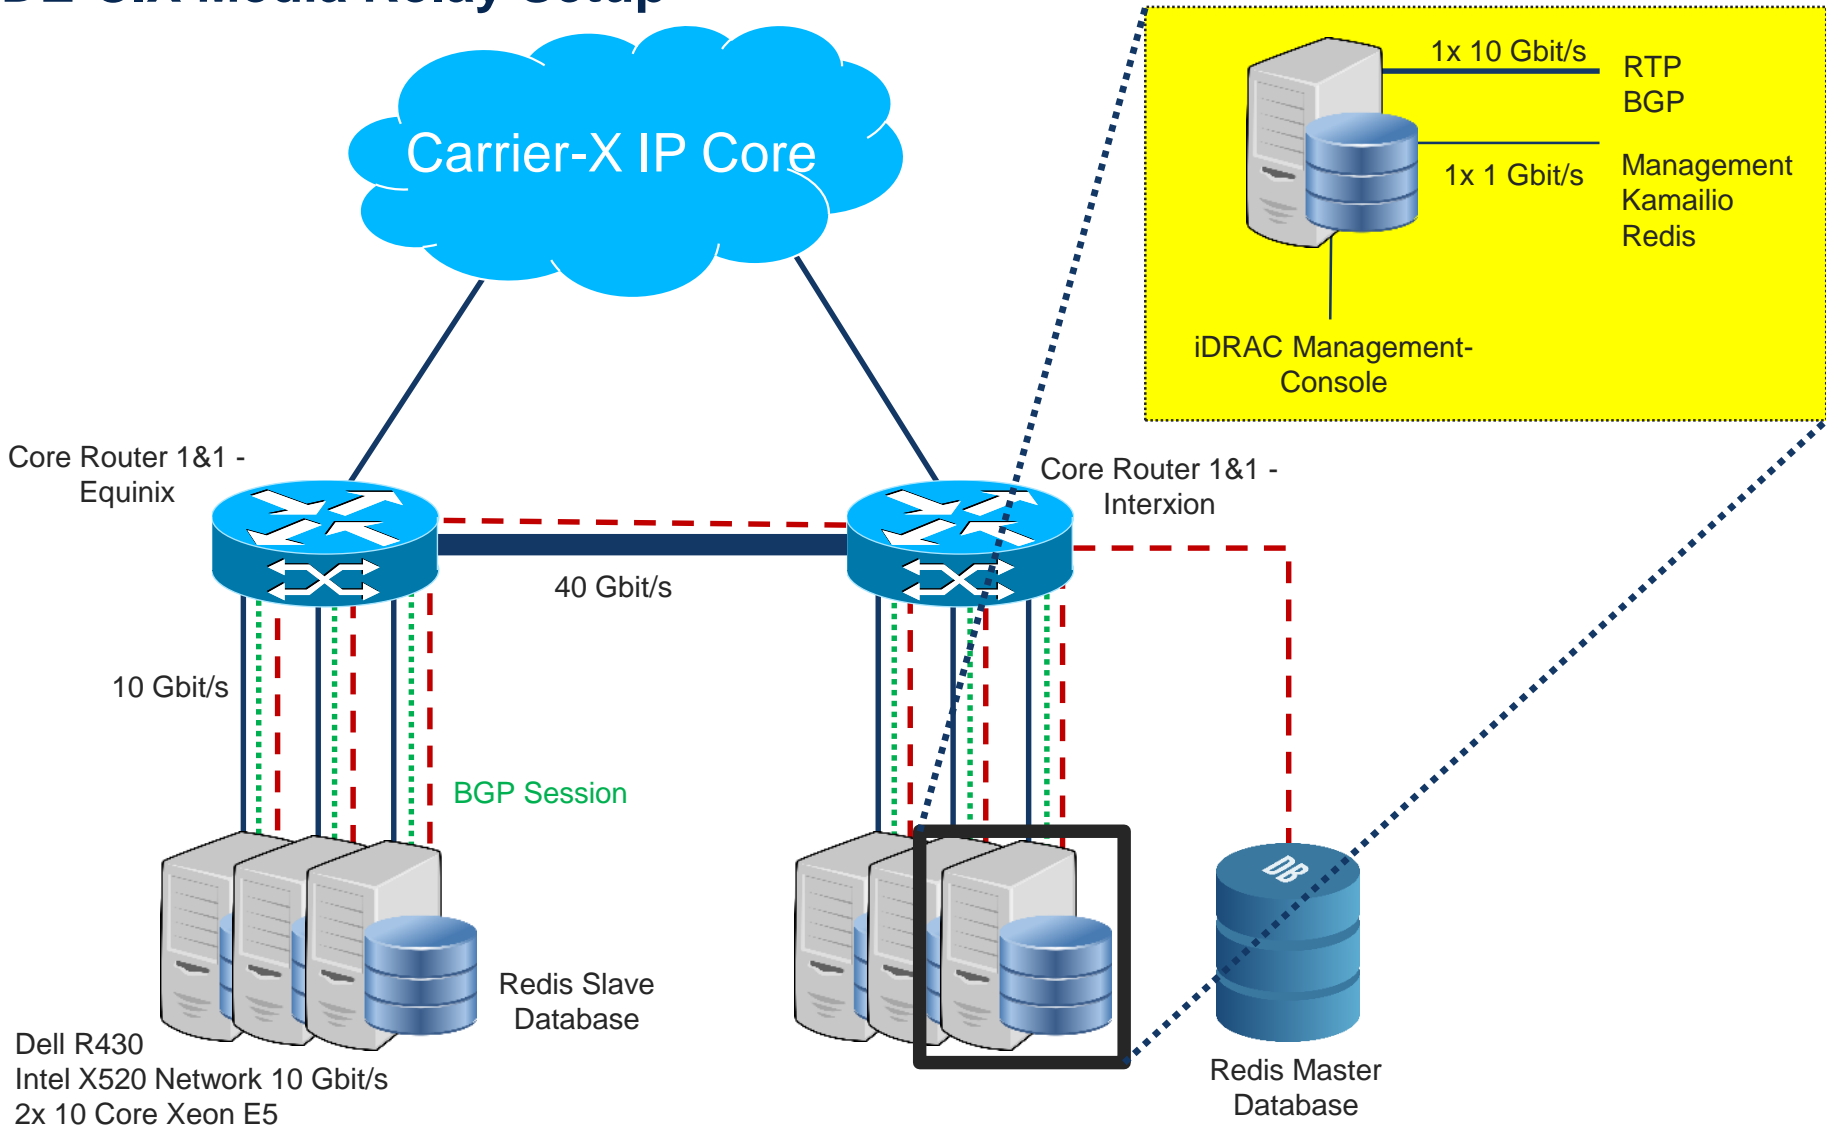
IRQ coalescing:
→ „Bundle“ interrupts




- Optimize utilization of CPU caches
- Improve performance
- Session mapped 1:1 to CPU core

Configured IRQ-CPU core affinity:
→ Map IRQ/queues to same CPU cores

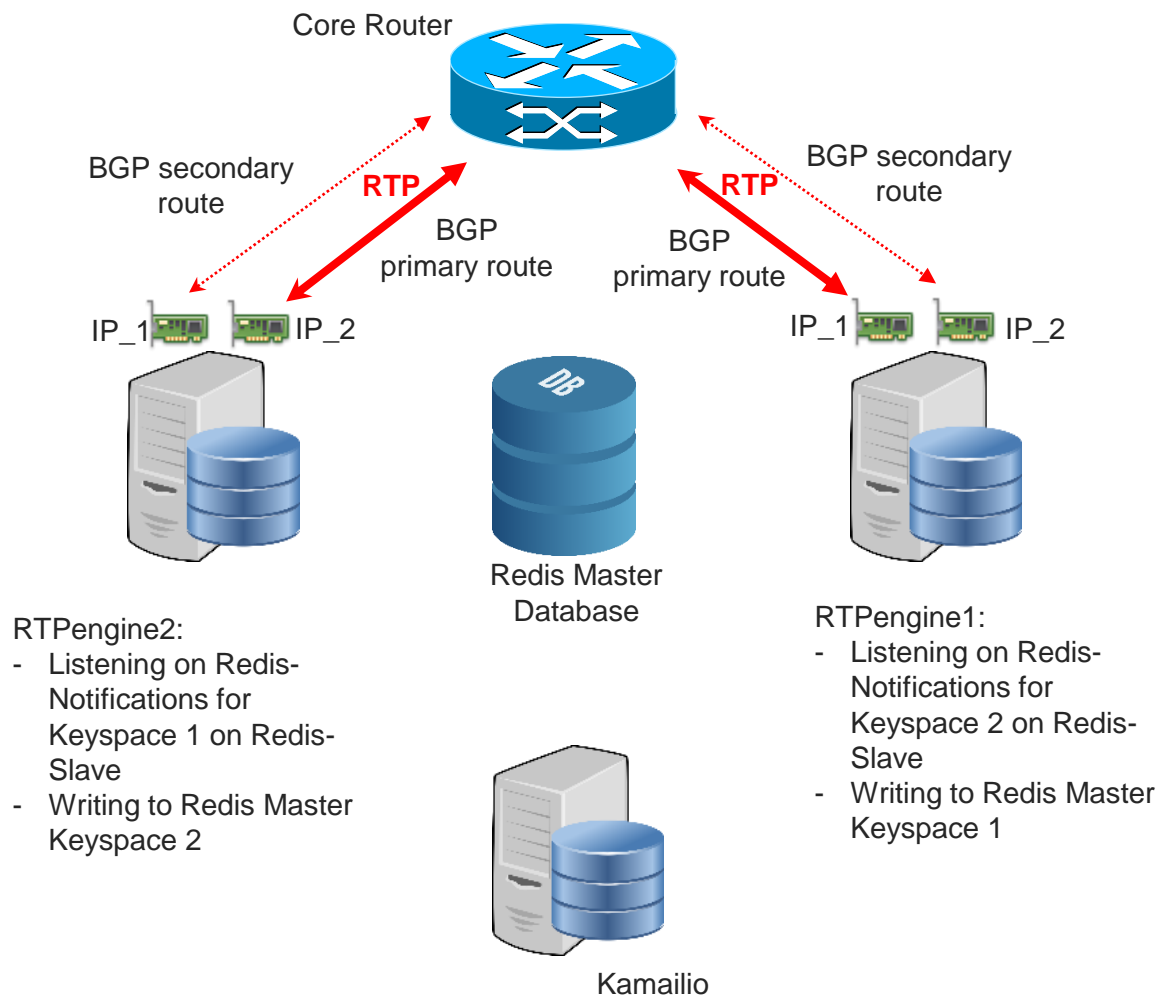
DE-CIX Media Relay Setup



RTPengine Redundancy

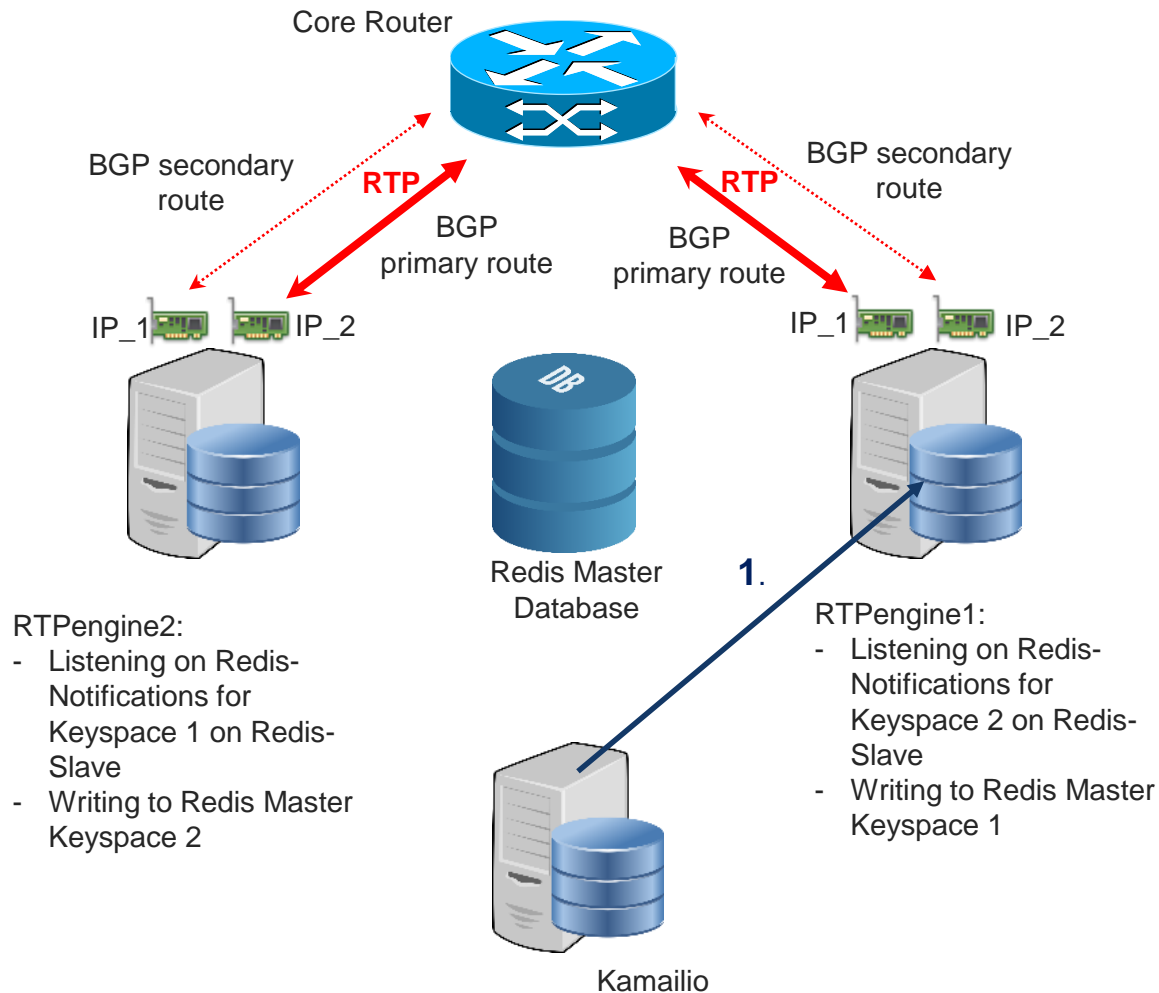


RTPengine Redundancy – Simplified illustration



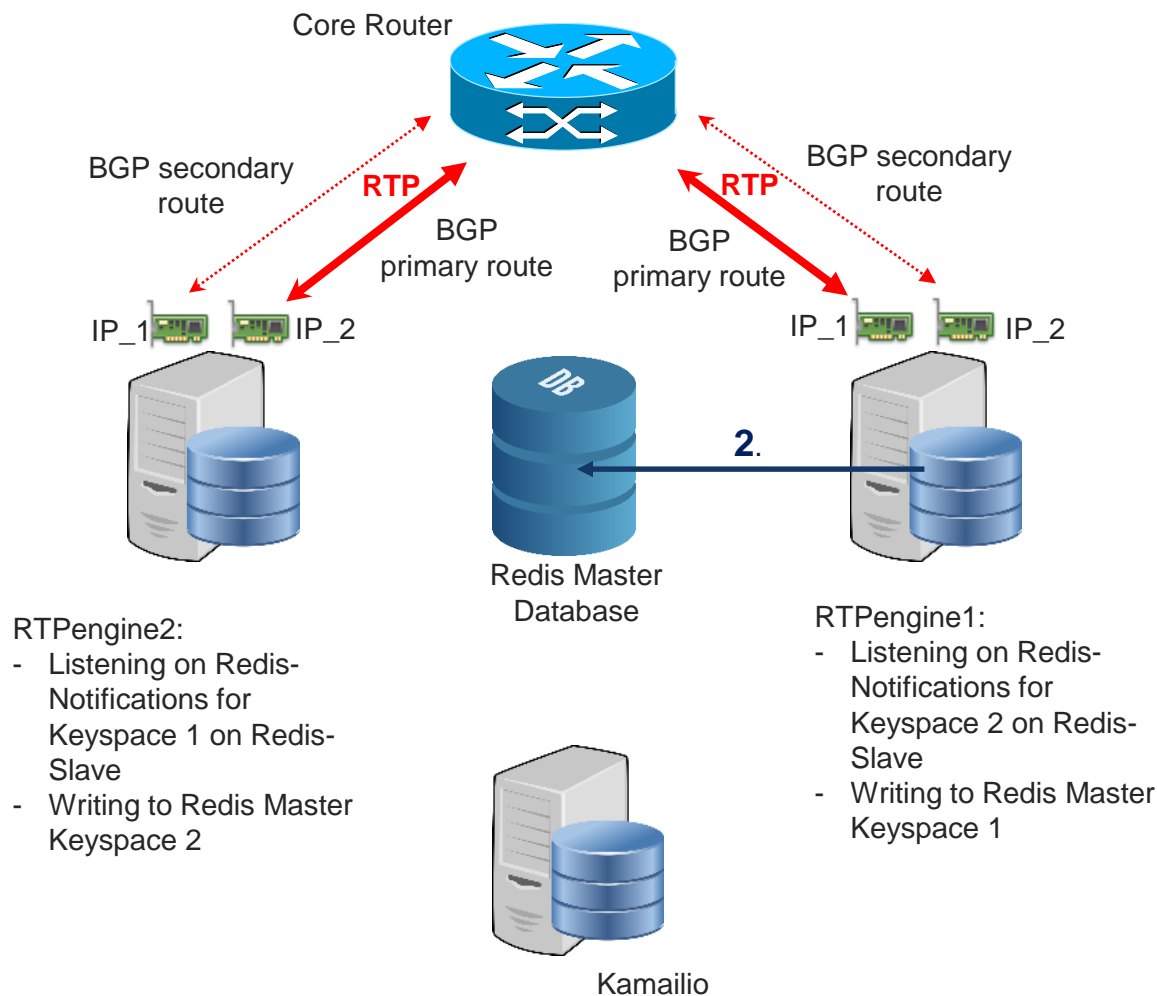
RTPengine Redundancy – Simplified illustration

1. Kamailio creates RTP session on RTPengine1 with active IP_1



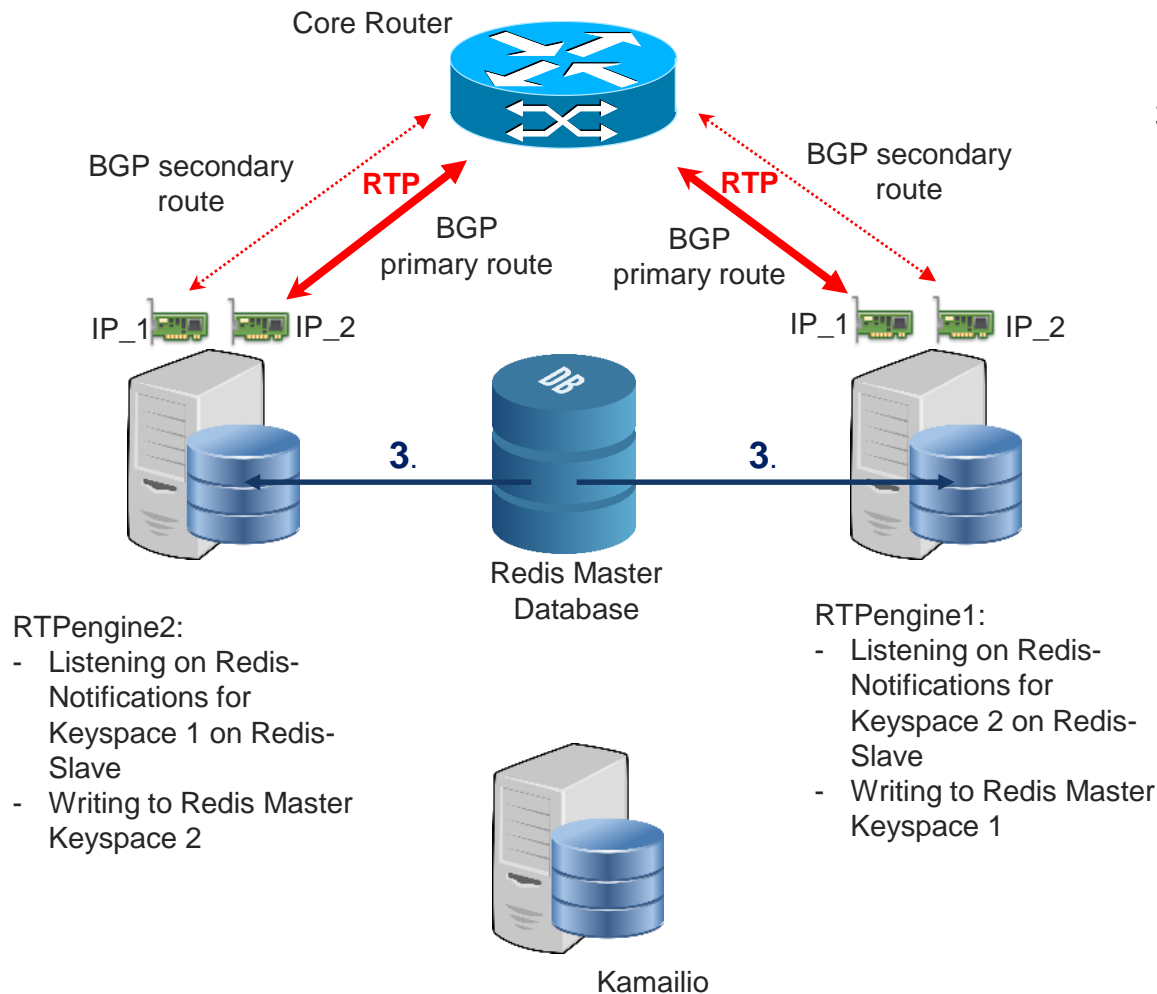
RTPengine Redundancy – Simplified illustration

1. Kamailio creates RTP session on RTPengine1 with active IP_1
2. **RTP traffic arrives. RTPengine1 persists session information to Redis Master, Keyspace 1**




RTPengine Redundancy – Simplified illustration

1. Kamailio creates RTP session on RTPengine1 with active IP_1
2. RTP traffic arrives. RTPengine1 persists session information to Redis Master, Keyspace 1
3. **Redis Master replicates information to slaves**

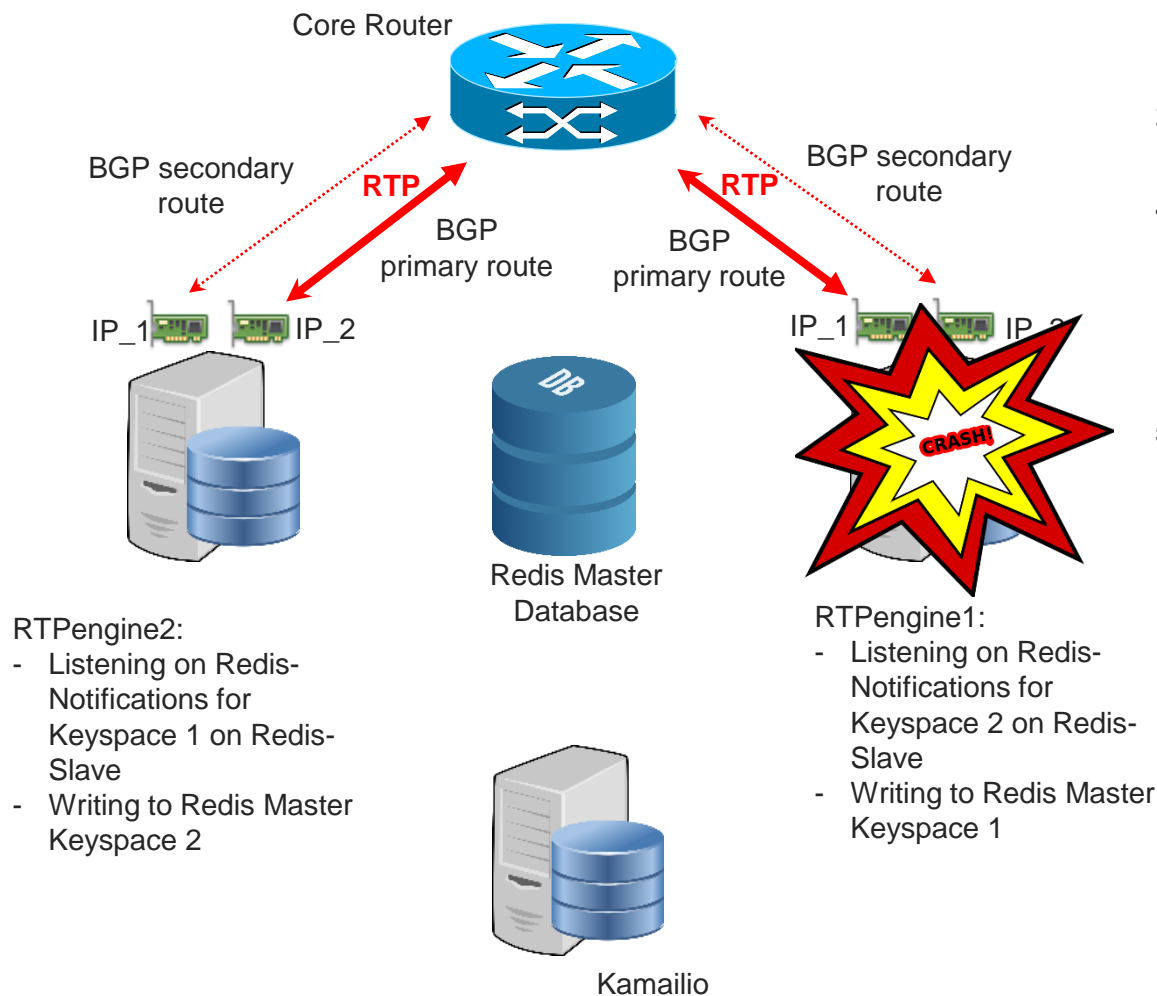


RTPEngine Redundancy – Simplified illustration

1. Kamailio creates RTP session on RTPEngine1 with active IP_1
2. RTP traffic arrives. RTPEngine1 persists session information to Redis Master, Keyspace 1
3. Redis Master replicates information to slaves
4. **RTPEngine2 (listening on Keyspace 1) handles Redis-Notification and creates session on IP_1**



RTPEngine Redundancy – Simplified illustration

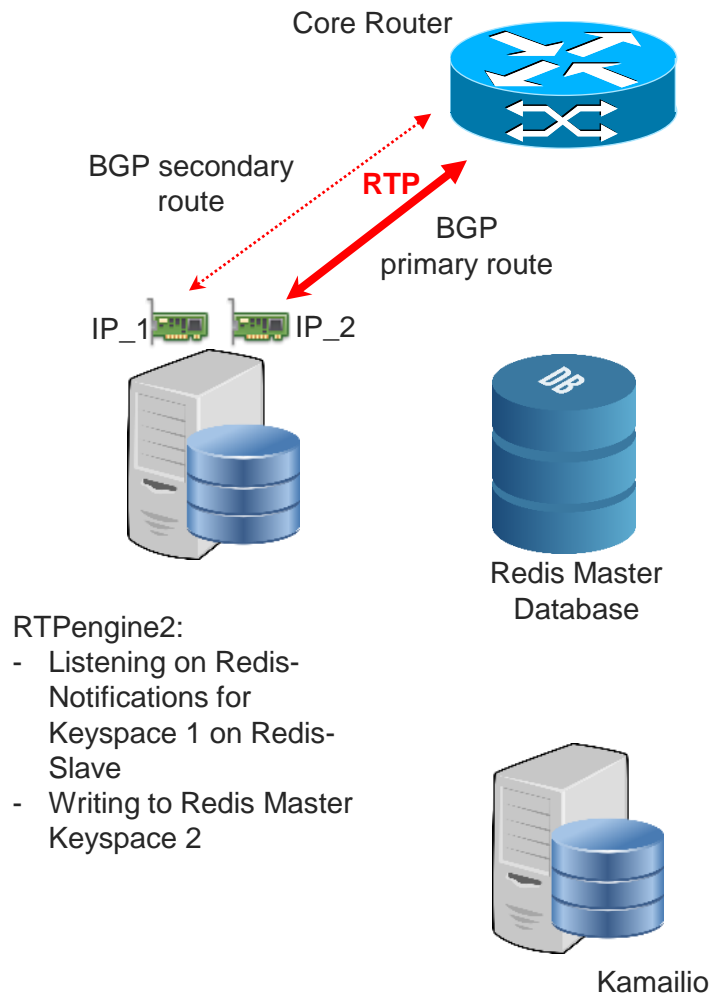


1. Kamailio creates RTP session on RTPEngine1 with active IP_1
2. RTP traffic arrives. RTPEngine1 persists session information to Redis Master, Keyspace 1
3. Redis Master replicates information to slaves
4. RTPEngine2 (listening on Keyspace 1) handles Redis-Notification and creates session on IP_1

5. HW Error / Network Error / ...




RTPEngine Redundancy – Simplified illustration


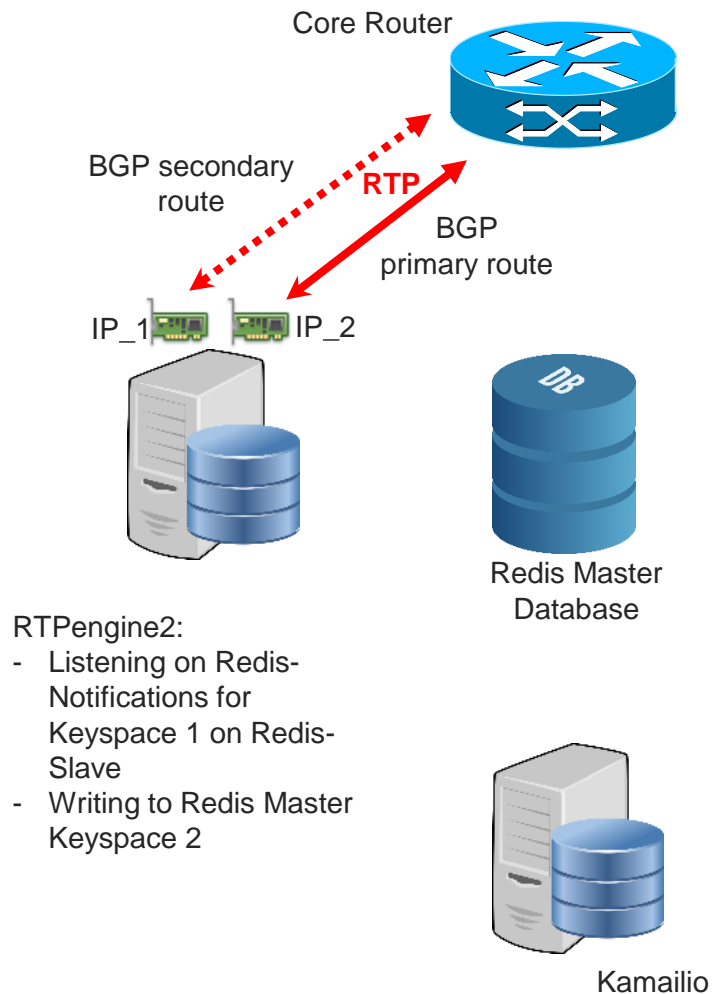


RTPEngine1:

- Listening on Redis-Notifications for Keyspace 2 on Redis-Slave
- Writing to Redis Master Keyspace 1

1. Kamailio creates RTP session on RTPEngine1 with active IP_1
2. RTP traffic arrives. RTPEngine1 persists session information to Redis Master, Keyspace 1
3. Redis Master replicates information to slaves
4. RTPEngine2 (listening on Keyspace 1) handles Redis-Notification and creates session on IP_1
5. HW Error / Network Error / ... 
6. **Core Router: Withdraw BGP announcements for RTPEngine1 (< 1s delay)**

RTPEngine Redundancy – Simplified illustration



RTPEngine1:

- Listening on Redis-Notifications for Keyspace 2 on Redis-Slave
- Writing to Redis Master Keyspace 1

1. Kamailio creates RTP session on RTPEngine1 with active IP_1
2. RTP traffic arrives. RTPEngine1 persists session information to Redis Master, Keyspace 1
3. Redis Master replicates information to slaves
4. RTPEngine2 (listening on Keyspace 1) handles Redis-Notification and creates session on IP_1

5. HW Error / Network Error / ...



6. Core Router: Withdraw BGP announcements for RTPEngine1 (< 1s delay)

7. Traffic routed via alternative route

Roundup

- Additional VoIP server in 2 new data centers
- Redundant active-active setup
- RTPengine software running on media relay servers
- Invoking RTPengine via Kamailio's rtpengine module
- Tuning IRQ handling
- Synchronizing RTPengine session information via Redis and keyspace notifications
- Using BGP mechanisms for redundancy



Interested?



➔ jobs.1und1.de